

```
In [1]: print("test")
```

test

```
In [101]: # First adding all necessary Libraries:
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder,StandardScaler
from sklearn.linear_model import LinearRegression,Lasso
from sklearn.metrics import mean_squared_error,mean_absolute_error
from sklearn.ensemble import RandomForestRegressor
import warnings
warnings.filterwarnings("ignore")
```

```
In [102]: # Loading the data of "LaptopPrice" dataset
df = pd.read_csv('D:laptopPrice.csv')
```

```
In [103]: # display the first five records
df.head(5)
```

Out[103]:

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	os_bit	graphic_card_gt
0	ASUS	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	64-bit	0 GE
1	Lenovo	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	64-bit	0 GE
2	Lenovo	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	64-bit	0 GE
3	ASUS	Intel	Core i5	10th	8 GB	DDR4	512 GB	0 GB	Windows	32-bit	2 GE
4	ASUS	Intel	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	Windows	64-bit	0 GE

```
In [104]: #The shape function is used to display the total number of rows and columns of the LaptopPrice dataset.
print(df.shape)
```

(823, 19)

```
In [105]: #Checking for null values in each column and displaying the sum of all null values in each column  
missing_values = df.isnull().sum()  
print("Missing Values:")  
print(missing_values)
```

```
Missing Values:  
brand                0  
processor_brand      0  
processor_name       0  
processor_gnrtn      0  
ram_gb               0  
ram_type             0  
ssd                  0  
hdd                  0  
os                   0  
os_bit              0  
graphic_card_gb     0  
weight              0  
warranty             0  
Touchscreen         0  
msoffice             0  
Price                0  
rating              0  
Number of Ratings   0  
Number of Reviews   0  
dtype: int64
```

```
In [106]: #Removing the rows with empty values
print(df.dropna())
```

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	\
0	ASUS	Intel	Core i3	10th	4 GB	DDR4	
1	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
2	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
3	ASUS	Intel	Core i5	10th	8 GB	DDR4	
4	ASUS	Intel	Celeron Dual	Not Available	4 GB	DDR4	
..	
818	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
819	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
820	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
821	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
822	Lenovo	AMD	Ryzen 5	10th	8 GB	DDR4	

	ssd	hdd	os	os_bit	graphic_card_gb	weight	\
0	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
1	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
2	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
3	512 GB	0 GB	Windows	32-bit	2 GB	Casual	
4	0 GB	512 GB	Windows	64-bit	0 GB	Casual	
..	
818	1024 GB	0 GB	Windows	64-bit	0 GB	Casual	
819	1024 GB	0 GB	Windows	64-bit	0 GB	Casual	
820	1024 GB	0 GB	Windows	64-bit	4 GB	Casual	
821	1024 GB	0 GB	Windows	64-bit	4 GB	Casual	
822	512 GB	0 GB	DOS	64-bit	0 GB	ThinNlight	

	warranty	Touchscreen	msoffice	Price	rating	Number of Ratings	\
0	No warranty	No	No	34649	2 stars	3	
1	No warranty	No	No	38999	3 stars	65	
2	No warranty	No	No	39999	3 stars	8	
3	No warranty	No	No	69990	3 stars	0	
4	No warranty	No	No	26990	3 stars	0	
..	
818	1 year	No	No	135990	3 stars	0	
819	1 year	No	No	144990	3 stars	0	
820	1 year	No	No	149990	3 stars	0	
821	1 year	No	No	142990	3 stars	0	
822	No warranty	No	No	57490	4 stars	18	

	Number of Reviews
0	0
1	5
2	1
3	0
4	0
..	...
818	0
819	0
820	0
821	0
822	4

[823 rows x 19 columns]

```
In [107]: # Checking if there is any duplicates value
df.duplicated().sum()
```

Out[107]: 21

```
In [108]: df = df.drop_duplicates()
```

```
In [109]: # Display basic information about the dataset
# info() is a method used to provide a summary of dataframe
# and understand the dataset .Also getting the structure of dataframe that am going to work on it.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 802 entries, 0 to 822
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   brand                  802 non-null    object
1   processor_brand        802 non-null    object
2   processor_name         802 non-null    object
3   processor_gnrtn        802 non-null    object
4   ram_gb                 802 non-null    object
5   ram_type               802 non-null    object
6   ssd                    802 non-null    object
7   hdd                    802 non-null    object
8   os                     802 non-null    object
9   os_bit                 802 non-null    object
10  graphic_card_gb        802 non-null    object
11  weight                 802 non-null    object
12  warranty               802 non-null    object
13  Touchscreen            802 non-null    object
14  msoffice                802 non-null    object
15  Price                  802 non-null    int64
16  rating                 802 non-null    object
17  Number of Ratings      802 non-null    int64
18  Number of Reviews      802 non-null    int64
dtypes: int64(3), object(16)
memory usage: 125.3+ KB
```

```
In [110]: #Checking the data types to see if all the data is in correct format.
# dtypes used to check and understanding the types of data presented in each column of the dataframe
df.dtypes
```

```
Out[110]: brand                  object
processor_brand                object
processor_name                  object
processor_gnrtn                 object
ram_gb                         object
ram_type                       object
ssd                            object
hdd                            object
os                             object
os_bit                         object
graphic_card_gb                object
weight                         object
warranty                       object
Touchscreen                    object
msoffice                       object
Price                          int64
rating                         object
Number of Ratings              int64
Number of Reviews              int64
dtype: object
```

```
In [111]: # count the number of unique values in each column of a DataFrame.
df.nunique()
```

```
Out[111]: brand                8
processor_brand              3
processor_name              11
processor_gnrtn             8
ram_gb                     4
ram_type                   6
ssd                        7
hdd                        4
os                         3
os_bit                    2
graphic_card_gb            5
weight                    3
warranty                   4
Touchscreen                2
msoffice                   2
Price                     405
rating                    5
Number of Ratings          282
Number of Reviews          135
dtype: int64
```

```
In [112]: #For numerical columns only
df.describe()
```

```
Out[112]:
```

	Price	Number of Ratings	Number of Reviews
count	802.000000	802.00000	802.000000
mean	76625.543641	299.84414	36.089776
std	45232.984422	1001.78442	118.313553
min	16990.000000	0.00000	0.000000
25%	45990.000000	0.00000	0.000000
50%	63990.000000	17.00000	2.000000
75%	89525.000000	140.25000	18.000000
max	441990.000000	15279.00000	1947.000000

```
In [113]: df.duplicated().sum()
```

```
Out[113]: 0
```

```
In [114]: # Checking the number of numeric features and cat_features (Categorical) from dataFram:
numeric_features = [feature for feature in df.columns if df[feature].dtype != 'object']
cat_features = [feature for feature in df.columns if df[feature].dtype == 'object']
# Display Numerical and Categorical variables
print(" Numerical features: ", numeric_features)
print("Categorical features:", cat_features)
```

```
Numerical features: ['Price', 'Number of Ratings', 'Number of Reviews']
Categorical features: ['brand', 'processor_brand', 'processor_name', 'processor_gnrtn', 'ram_gb', 'ram_
type', 'ssd', 'hdd', 'os', 'os_bit', 'graphic_card_gb', 'weight', 'warranty', 'Touchscreen', 'msoffic
e', 'rating']
```

```
In [115]: # describe () method used to describe numerical column only
df.describe(include = 'object')
```

Out[115]:

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	os_bit	graphic_card
count	802	802	802	802	802	802	802	802	802	802	
unique	8	3	11	8	4	6	7	4	3	2	
top	ASUS	Intel	Core i5	11th	8 GB	DDR4	512 GB	0 GB	Windows	64-bit	
freq	243	594	284	328	404	690	389	602	763	693	

```
In [117]: df.loc[df['Price'] == 1, 'Price'] = 500
```

```
In [118]: df.describe()
```

Out[118]:

	Price	Number of Ratings	Number of Reviews
count	802.000000	802.00000	802.000000
mean	76625.543641	299.84414	36.089776
std	45232.984422	1001.78442	118.313553
min	16990.000000	0.00000	0.000000
25%	45990.000000	0.00000	0.000000
50%	63990.000000	17.00000	2.000000
75%	89525.000000	140.25000	18.000000
max	441990.000000	15279.00000	1947.000000

```
In [119]: df['Price'].describe()
```

Out[119]: count 802.000000
mean 76625.543641
std 45232.984422
min 16990.000000
25% 45990.000000
50% 63990.000000
75% 89525.000000
max 441990.000000
Name: Price, dtype: float64

```
In [120]: df.describe(include = 'object')#summary statistics for categorical values
```

Out[120]:


	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	os_bit	graphic_card
count	802	802	802	802	802	802	802	802	802	802	
unique	8	3	11	8	4	6	7	4	3	2	
top	ASUS	Intel	Core i5	11th	8 GB	DDR4	512 GB	0 GB	Windows	64-bit	
freq	243	594	284	328	404	690	389	602	763	693	

```
In [121]: numeric_features = [feature for feature in df.columns if df[feature].dtype != 'object']
cat_features = [feature for feature in df.columns if df[feature].dtype == 'object']
print("Numerical features: ", numeric_features)
print("Categorical features:", cat_features)

Numerical features: ['Price', 'Number of Ratings', 'Number of Reviews']
Categorical features: ['brand', 'processor_brand', 'processor_name', 'processor_gnrtn', 'ram_gb', 'ram_type', 'ssd', 'hdd', 'os', 'os_bit', 'graphic_card_gb', 'weight', 'warranty', 'Touchscreen', 'msoffice', 'rating']
```

```
In [122]: import seaborn as sns
plt.figure(figsize=(10,6))
sns.regplot(x="Number of Ratings", y="Price", data=df)
```

Out[122]: <matplotlib.axes._subplots.AxesSubplot at 0x122d257adc0>




```
In [123]: from scipy import stats
pearson_coef, p_value = stats.pearsonr(df['Number of Ratings'], df['Price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is -0.15255276430421938 with a P-value of P = 1.4318727176497412e-05

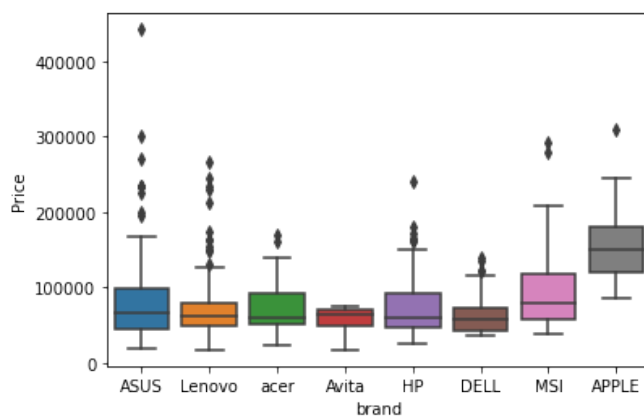
```
In [124]: plt.figure(figsize=(10,6))  
sns.regplot(x="Number of Reviews", y="Price", data=df)
```

```
Out[124]: <matplotlib.axes._subplots.AxesSubplot at 0x122d26d7250>
```




```
In [125]: # In the given plot below, it is observed that the price range vary for ASUS and Apple Brand.  
# This indicates the categories can vary with price hence features can be used for prediction.  
sns.boxplot(x="brand", y="Price", data=df)
```

```
Out[125]: <matplotlib.axes._subplots.AxesSubplot at 0x122d27504f0>
```





```
In [126]: plt.figure(figsize=(10,6))
sns.boxplot(x="processor_brand", y="Price", data=df)
```

Out[126]: <matplotlib.axes._subplots.AxesSubplot at 0x122d29bfa90>




```
In [127]: plt.figure(figsize=(10,6))
sns.boxplot(x="processor_name", y="Price", data=df)
```

Out[127]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2a42250>



```
In [128]: plt.figure(figsize=(10,6))  
sns.boxplot(x="processor_gnrtn", y="Price", data=df)
```


Out[128]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2ca36d0>



```
In [ ]: # processor_name feature shows a huge difference in price ranges between laptop processor name with Core  
# This feature is very important for price prediction as the bigger the difference in range the better
```


```
In [129]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ram_gb", y="Price", data=df)
```

Out[129]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2ca3610>




```
In [130]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ram_type", y="Price", data=df)
```

Out[130]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2b5c940>




```
In [131]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ssd", y="Price", data=df)
```

Out[131]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2d6a6d0>




```
In [132]: plt.figure(figsize=(10,6))  
sns.boxplot(x="hdd", y="Price", data=df)
```

Out[132]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2747220>




```
In [133]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os", y="Price", data=df)
```

Out[133]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2d91700>




```
In [134]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os_bit", y="Price", data=df)
```

Out[134]: <matplotlib.axes._subplots.AxesSubplot at 0x122d26b00d0>




```
In [135]: plt.figure(figsize=(10,6))  
sns.boxplot(x="graphic_card_gb", y="Price", data=df)
```

Out[135]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2db06d0>




```
In [136]: plt.figure(figsize=(10,6))  
sns.boxplot(x="weight", y="Price", data=df)
```

```
Out[136]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2e6b490>
```




```
In [137]: plt.figure(figsize=(10,6))  
sns.boxplot(x="warranty", y="Price", data=df)
```

```
Out[137]: <matplotlib.axes._subplots.AxesSubplot at 0x122d2ed4b20>
```




```
In [138]: plt.figure(figsize=(10,6))  
sns.boxplot(x="Touchscreen", y="Price", data=df)
```

```
Out[138]: <matplotlib.axes._subplots.AxesSubplot at 0x122d3f41f40>
```




```
In [139]: plt.figure(figsize=(10,6))  
sns.boxplot(x="msoffice", y="Price", data=df)
```

```
Out[139]: <matplotlib.axes._subplots.AxesSubplot at 0x122d3fb1f10>
```




```
In [140]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os_bit", y="Price", data=df)
```

```
Out[140]: <matplotlib.axes._subplots.AxesSubplot at 0x122d4013ca0>
```



```
In [141]: plt.figure(figsize=(10,6))  
sns.boxplot(x="rating", y="Price", data=df)
```

```
Out[141]: <matplotlib.axes._subplots.AxesSubplot at 0x122d40afeb0>
```



```
In [142]: df.drop(['weight', 'warranty', 'Touchscreen', 'processor_brand', 'os_bit'], axis = 1, inplace = True)
```



```
In [143]: df
```

Out[143]:

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic_card_gb	msoffice	Price	ra
0	ASUS	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	34649	;
1	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	38999	;
2	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	39999	;
3	ASUS	Core i5	10th	8 GB	DDR4	512 GB	0 GB	Windows	2 GB	No	69990	;
4	ASUS	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	Windows	0 GB	No	26990	;
...	;
818	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	0 GB	No	135990	;
819	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	0 GB	No	144990	;
820	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	4 GB	No	149990	;
821	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	4 GB	No	142990	;
822	Lenovo	Ryzen 5	10th	8 GB	DDR4	512 GB	0 GB	DOS	0 GB	No	57490	;

802 rows x 14 columns

```
In [144]: df
```

Out[144]:

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic_card_gb	msoffice	Price	ra
0	ASUS	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	34649	;
1	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	38999	;
2	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	0 GB	No	39999	;
3	ASUS	Core i5	10th	8 GB	DDR4	512 GB	0 GB	Windows	2 GB	No	69990	;
4	ASUS	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	Windows	0 GB	No	26990	;
...	;
818	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	0 GB	No	135990	;
819	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	0 GB	No	144990	;
820	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	4 GB	No	149990	;
821	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	4 GB	No	142990	;
822	Lenovo	Ryzen 5	10th	8 GB	DDR4	512 GB	0 GB	DOS	0 GB	No	57490	;

802 rows x 14 columns

```
In [145]: #brand', 'processor_name', 'processor_gnrtn', 'ram_gb', 'ram_type', 'ssd', 'hdd', 'os',
#graphic_card_gb', 'msoffice', 'rating'
from sklearn.preprocessing import LabelEncoder

labelencoder = LabelEncoder()
df.brand = labelencoder.fit_transform(df.brand)
df.processor_name = labelencoder.fit_transform(df.processor_name)
df.processor_gnrtn = labelencoder.fit_transform(df.processor_gnrtn)
df.ram_gb = labelencoder.fit_transform(df.ram_gb)
df.ram_type = labelencoder.fit_transform(df.ram_type)
df.ssd = labelencoder.fit_transform(df.ssd)
df.hdd = labelencoder.fit_transform(df.hdd)
df.os = labelencoder.fit_transform(df.os)
df.graphic_card_gb = labelencoder.fit_transform(df.graphic_card_gb)
df.msoffice = labelencoder.fit_transform(df.msoffice)
df.rating = labelencoder.fit_transform(df.rating)
from sklearn.preprocessing import LabelEncoder
```

```
In [168]: import scipy.stats as stats
df = stats.zscore(df)
```

```
In [170]: df
```

```
Out[170]: array([[ -1.15409599, -0.87301913, -0.91832804, ..., -2.76166878,
-0.2965003 , -0.30522536],
 [ 0.87903961, -0.87301913, -0.91832804, ..., -1.00105005,
-0.23457211, -0.2629384 ],
 [ 0.87903961, -0.87301913, -0.91832804, ..., -1.00105005,
-0.29150609, -0.29676797],
 ...,
 [ -1.15409599,  2.20686947,  1.48628652, ..., -1.00105005,
-0.29949682, -0.30522536],
 [ -1.15409599,  2.20686947,  1.48628652, ..., -1.00105005,
-0.29949682, -0.30522536],
 [ 0.87903961,  1.52244978, -0.91832804, ...,  0.75956868,
-0.28151767, -0.27139579]])
```

```
In [171]: x_train=df.iloc[:,0:13]
y_train=df.iloc[:,10]
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-171-e365df6435b8> in <module>
----> 1 x_train=df.iloc[:,0:13]
      2 y_train=df.iloc[:,10]

AttributeError: 'numpy.ndarray' object has no attribute 'iloc'
```

```
In [172]: x_train.head()
```

Out[172]:

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic_card_gb	msoffice	Price	rating	Number of Reviews
0	1	1	0	2	1	0	1	2	0	0	34649	1	6
1	5	1	0	2	1	0	1	2	0	0	38999	2	6
2	5	1	0	2	1	0	1	2	0	0	39999	2	6
3	1	2	0	3	1	6	0	2	1	0	69990	2	6
4	1	0	7	2	1	0	3	2	0	0	26990	2	6

```
In [173]: y_train.head()
```

```
Out[173]: 0    34649
          1    38999
          2    39999
          3    69990
          4    26990
          Name: Price, dtype: int64
```

```
In [174]: # importing train_test_split from sklearn
          from sklearn.model_selection import train_test_split
          # splitting the data # 30% for testing is used
          X_train, X_test, Y_train, Y_test = train_test_split(x_train, y_train, test_size = 0.3, random_state = 0)
```

```
In [175]: #Multiple Linear Regression
          from sklearn.linear_model import LinearRegression

          model = LinearRegression()
          model_mlr = model.fit(X_train,Y_train)
```

```
In [176]: #Making price prediction using the testing set (Fit to MLR)
          Y_pred_MLR = model_mlr.predict(X_test)
```

```
In [177]: #Calculating the Mean Square Error for MLR model
          mse_MLR = mean_squared_error(Y_test, Y_pred_MLR)
          print('The mean square error for Multiple Linear Regression: ', mse_MLR)

          The mean square error for Multiple Linear Regression:  3.4493088314991315e-22
```

```
In [178]: #The mean square error for Multiple Linear Regression:  0.3674647167443785
```

```
In [179]: #Calculating the Mean Absolute Error for MLR model
          mae_MLR= mean_absolute_error(Y_test, Y_pred_MLR)
          print('The mean absolute error for Multiple Linear Regression: ', mae_MLR)

          The mean absolute error for Multiple Linear Regression:  1.5533112831939305e-11
```

```
In [180]: #Calling the random forest model and fitting the training data
          rfModel = RandomForestRegressor()
          model_rf = rfModel.fit(X_train,Y_train)
```

```
In [181]: #Prediction of Laptop prices using the testing data
          Y_pred_RF = model_rf.predict(X_test)
```

```
In [182]: #Calculating the Mean Square Error for Random Forest Model
          mse_RF = mean_squared_error(Y_test, Y_pred_RF)
          print('The mean square error of price and predicted value is: ', mse_RF)

          The mean square error of price and predicted value is:  2941640.472541086
```

```
In [183]: #Calculating the Mean Absolute Error for Random Forest Model
          mae_RF= mean_absolute_error(Y_test, Y_pred_RF)
          print('The mean absolute error of price and predicted value is: ', mae_RF)

          The mean absolute error of price and predicted value is:  317.3755186721992
```

```
In [184]: #LASSO Model
          #Calling the model and fitting the training data
          LassoModel = Lasso()
          model_lm = LassoModel.fit(X_train,Y_train)
```

```
In [185]: #Price prediction using testing data
          Y_pred_lasso = model_lm.predict(X_test)
```

```
In [186]: #Mean Absolute Error for LASSO Model
mae_lasso= mean_absolute_error(Y_test, Y_pred_lasso)
print('The mean absolute error of price and predicted value is: ', mae_lasso)
```

The mean absolute error of price and predicted value is: 1.3399300380673692e-05

```
In [187]: #Mean Squared Error for the LASSO Model
mse_lasso = mean_squared_error(Y_test, Y_pred_lasso)
print('The mean square error of price and predicted value is: ', mse_lasso)
```

The mean square error of price and predicted value is: 3.446638548304528e-10

```
In [188]: scores = [('MLR', mae_MLR),
                    ('Random Forest', mae_RF),
                    ('LASSO', mae_lasso)
                  ]
```

```
In [189]: mae = pd.DataFrame(data = scores, columns=['Model', 'MAE Score'])
mae
```

Out[189]:


	Model	MAE Score
0	MLR	1.553311e-11
1	Random Forest	3.173755e+02
2	LASSO	1.339930e-05

```
In [ ]: # By observing MAE score of MLR, Random Forest and LASSO I observe that random Forest algorithm has the
# in general: the lower MAE score the better model.
# Hence , I conclude that Random Forest is the better model among the selected three models
```

```
In [190]: mae.sort_values(by=(['MAE Score']), ascending=False, inplace=True)

f, axe = plt.subplots(1,1, figsize=(10,7))
sns.barplot(x = mae['Model'], y=mae['MAE Score'], ax = axe)
axe.set_xlabel('Model', size=20)
axe.set_ylabel('Mean Absolute Error', size=20)

plt.show()
```



In []:

In []: